Equivariant Flow Matching for Generative Cardiac Shape Modelling

Sean Jaffe
TODO: affiliation

March 10, 2026

Abstract

We present an SE(3)-equivariant flow matching model for generative cardiac shape modelling directly on surface meshes. Given registered left-ventricular meshes extracted from the ACDC cine-MRI dataset, our model learns a conditional velocity field over mesh vertex positions using an $E(n)$ Equivariant Graph Neural Network (EGNN), trained with the optimal-transport conditional flow matching objective. A VAE-style encoder compresses each mesh into a 64-dimensional latent code that captures pathology-relevant shape variation. On reconstruction quality, our model achieves a Chamfer distance of 0.968 ± 0.078 mm and Hausdorff distance of 1.576 ± 0.094 mm, an $8.7\times$ improvement in Chamfer distance over the PCA shape model baseline (8.40 ± 1.04 mm). In distribution quality, 1-NNA of 0.600 indicates generated shapes are partially indistinguishable from real data, and a linear probe on the latent space achieves 33% accuracy across five unsupervised pathology classes (chance: 20%), demonstrating that the latent space encodes clinically meaningful structure without label supervision.

1 Introduction

[TODO: 3–4 paragraphs covering:]

- Clinical motivation: cardiac shape is a biomarker for disease (DCM, HCM, myocardial infarction).
- Gap: existing statistical shape models (PCA) are linear and voxel diffusion models discard surface topology.
- Our contribution: SE(3)-equivariant GNN combined with flow matching directly on cardiac meshes.
- Summary of results.

2 Background

2.1 Statistical Shape Models

Statistical shape models (SSMs) represent shape variation within a population by applying principal component analysis (PCA) to a set of registered point correspondences Cootes et al. [1995]. Given N shapes with V vertices each, we flatten each shape to a vector $\mathbf{s}_i \in \mathbb{R}^{3V}$ and compute the empirical covariance. New shapes are sampled as

$$\hat{\mathbf{s}} = \bar{\mathbf{s}} + \mathbf{P}\mathbf{b}, \quad \mathbf{b} \sim \mathcal{N}(\mathbf{0}, \mathbf{\Lambda}), \tag{1}$$

where $\bar{\mathbf{s}}$ is the mean shape, $\mathbf{P}$ are the principal modes, and $\mathbf{\Lambda}$ is the diagonal matrix of eigenvalues. SSMs are computationally efficient but capture only linear variation, failing to represent non-linear deformation modes such as the asymmetric remodelling seen in hypertrophic cardiomyopathy.

2.2 Geometric Deep Learning and SE(3) Equivariance

A function f is *equivariant* under a symmetry group G if

$$f(g \cdot \mathbf{x}) = g \cdot f(\mathbf{x}) \quad \forall\, g \in G,\ \mathbf{x} \in \mathcal{X}. \tag{2}$$

For 3-D point sets, the relevant group is $\mathrm{SE}(3) = \mathrm{SO}(3) \ltimes \mathbb{R}^3$, the group of rotations and translations in three dimensions. A neural network operating on point positions that is equivariant to $\mathrm{SE}(3)$ will produce geometrically consistent outputs regardless of the coordinate frame in which the input is expressed.

We use the $E(n)$ *Equivariant Graph Neural Network* (EGNN) of Satorras et al. [2021] as our equivariant backbone. Each EGNN layer jointly updates scalar node features $\mathbf{h}_i$ and equivariant coordinates $\mathbf{x}_i$:

$$\mathbf{m}_{ij} = \phi_e\Big(\mathbf{h}_i,\, \mathbf{h}_j,\, \|\mathbf{x}_i - \mathbf{x}_j\|^2\Big)\,, \tag{3}$$

$$\mathbf{x}_i \leftarrow \mathbf{x}_i + \sum_{j \neq i} (\mathbf{x}_i - \mathbf{x}_j)\, \phi_x(\mathbf{m}_{ij}), \tag{4}$$

$$\mathbf{h}_i \leftarrow \phi_h\left(\mathbf{h}_i,\, \sum_j \mathbf{m}_{ij}\right)\,, \tag{5}$$

where ϕ_e, ϕ_x, and ϕ_h are multi-layer perceptrons. Because the coordinate update in (4) depends only on relative positions and the scalar ϕ_x, the overall mapping $(\mathbf{H}, \mathbf{X}) \mapsto (\mathbf{H}', \mathbf{X}')$ is equivariant to $\mathrm{SE}(3)$ by construction.

2.3 Diffusion Models and Flow Matching

Score-based diffusion models Song et al. [2021] learn to reverse a forward noising process, but typically require thousands of denoising steps at inference time. *Flow matching* Lipman et al. [2023] instead defines a deterministic probability path between a noise distribution p_0 and the data distribution p_1, parameterising the corresponding velocity field $v_\theta(\mathbf{x}_t, t)$ that transports samples from p_0 to p_1.

Under the *Optimal Transport* (OT) path Albergo and Vanden-Eijnden [2023], interpolating trajectories are straight lines:

$$\mathbf{x}_t = (1 - t)\, \mathbf{x}_0 + t\, \mathbf{x}_1, \quad t \sim \mathcal{U}[0, 1], \tag{6}$$

yielding the simple regression objective

$$\mathcal{L}_{\text{CFM}}(\theta) = \mathbb{E}_{t,\, \mathbf{x}_0 \sim p_0,\, \mathbf{x}_1 \sim p_1}\left[\|v_\theta(\mathbf{x}_t, t, \mathbf{z}) - (\mathbf{x}_1 - \mathbf{x}_0)\|^2\right], \tag{7}$$

where $\mathbf{z}$ is an optional conditioning variable. Samples are generated by integrating v_θ from $t = 0$ to $t = 1$ using a standard ODE solver (Euler or RK4). Compared to score-based methods, flow matching is faster to train (single forward pass per sample) and requires far fewer integration steps at inference.

3 Method

3.1 Data and Preprocessing

We use the *Automated Cardiac Diagnosis Challenge* (ACDC) dataset Bernard et al. [2018], which provides cine-MRI volumes with expert segmentation masks for 100 subjects spanning five pathology classes: normal (NOR), myocardial infarction (MINF), dilated cardiomyopathy (DCM), hypertrophic cardiomyopathy (HCM), and right ventricular abnormality (RV). Segmentations label the left ventricle (LV, label 3), right ventricle (RV cavity, label 1), and myocardium (Myo, label 2) at end-diastole (ED) and end-systole (ES).

Mesh extraction. For each structure at each frame, we apply the marching cubes algorithm Lorensen and Cline [1987] to the binary segmentation mask, using the physical voxel spacing as the step size to produce meshes in millimetre coordinates. We then apply ten iterations of Laplacian smoothing to reduce staircase artefacts, and decimate each mesh to approximately $2{,}500$ vertices via quadric edge collapse, targeting a balance between geometric fidelity and computational cost.

figures/ed_lv_single.pdf

(a) ED LV mesh after marching cubes, Laplacian smoothing, and quadric decimation ($\sim$2,500 vertices).

figures/ed_es_comparison.pdf

(b) Same subject at end-diastole (blue, larger) and end-systole (red, contracted). Volume difference gives the ejection fraction.

Figure 1: Representative LV surface meshes extracted from ACDC segmentation masks.

Template registration. To establish vertex correspondence across subjects we select a single reference subject and register all others to it. Registration proceeds in two stages: (i) a coarse Procrustes alignment using the Kabsch algorithm, and (ii) refinement with Iterative Closest Point (ICP). After alignment, each mesh is centred at the origin by subtracting its centroid, removing translation as a degree of freedom.

Graph construction. Each registered mesh is stored as a PyTorch Geometric `Data` object containing vertex positions $\mathbf{X} \in \mathbb{R}^{V \times 3}$, a k-nearest-neighbour edge index ($k = 8$) computed in the registered coordinate frame, the face connectivity for surface-area and volume calculations, and an integer pathology label $y \in \{0, \ldots, 4\}$.

3.2 Model Architecture

Our model consists of two learned components: a graph encoder that maps a cardiac mesh to a latent code, and a flow-matching velocity field conditioned on that code. Figure 2 provides an overview.

[TODO: Insert architecture diagram: mesh → encoder → z → velocity field → generated mesh]

Figure 2: Model architecture. The EGNN encoder compresses a cardiac mesh graph into a latent code $\mathbf{z}$. The EGNN velocity field, conditioned on $\mathbf{z}$ and time t, transports a noise cloud to a realistic cardiac mesh via OT flow matching.

EGNN Encoder. The encoder is a four-layer EGNN followed by mean-and-max global pooling, producing a Gaussian posterior $q_\phi(\mathbf{z} \mid \mathcal{G}) = \mathcal{N}(\boldsymbol{\mu}, \mathrm{diag}(\boldsymbol{\sigma}^2))$ with $\boldsymbol{\mu}, \log \boldsymbol{\sigma} \in \mathbb{R}^{64}$. Initial node features are set to a constant vector (no semantic node attributes are available), so all information must be inferred from the geometric structure of the graph.

Equivariant Velocity Field. The velocity field $v_\theta(\mathbf{x}_t, t, \mathbf{z})$ is a four-layer EGNN where each node's input feature is the concatenation of a sinusoidal time embedding and the conditioning code $\mathbf{z}$ projected to the hidden dimension. The network outputs a velocity vector in $\mathbb{R}^3$ for each vertex. Because position updates in EGNN depend only on relative coordinates, the velocity field inherits $\mathrm{SE}(3)$-equivariance: rotating the input mesh rotates the predicted velocities by the same rotation.

Training Objective. The full objective is

$$\mathcal{L} = \mathcal{L}_{\text{CFM}}(\theta) + \beta \, D_{\text{KL}}(q_\phi(\mathbf{z} \mid \mathcal{G}) \parallel \mathcal{N}(\mathbf{0}, \mathbf{I})) \,, \tag{8}$$

where $\mathcal{L}_{\text{CFM}}$ is the OT flow matching loss from Equation (7) and β is a KL weight annealed from 0 to 5×10^{-3} over the first 50 epochs.

3.3 Baselines

We compare against four baselines, each designed to ablate a specific modelling choice.

Baseline 1: PCA Shape Model. Classical linear SSM fit to flattened vertex positions. Justifies the use of deep learning at all.

Baseline 2: Voxel Diffusion. OT flow matching applied to binary segmentation volumes on a 32^3 grid, using a 3-D U-Net velocity field. Justifies the mesh representation over voxel grids.

Baseline 3: PointNet++ VAE. PointNet++ encoder with a FoldingNet decoder trained as a VAE (no diffusion/flow). Justifies the distributional modelling component.

Baseline 4: Non-equivariant Flow Matching. Identical architecture to our main model, but EGNN layers replaced by standard MLP message passing; random rotation augmentation is applied during training to compensate. This is the *critical* ablation: it directly validates the geometric deep learning contribution.

4 Experiments

4.1 Setup

All experiments use the ACDC dataset (100 subjects, 5 pathology classes). We filter to end-diastole left-ventricular meshes and use an 80/20 train/validation split. The main model is trained for 200

epochs with AdamW (lr $= 3 \times 10^{-4}$, weight decay $= 10^{-4}$), batch size 8, cosine annealing, and gradient clipping at 1.0. The encoder has 4 EGNN layers; the flow model has 6 EGNN layers; hidden dimension is 128 and latent dimension 64. KL weight $\beta = 5 \times 10^{-3}$ with linear warmup over the first 50 epochs (from $\beta = 0$ to $\beta = 5 \times 10^{-3}$). These hyperparameters were selected via the ablation study described in Section 4.6. Training runs on a single NVIDIA GPU. Full training curves and hyperparameter sweeps are logged to the W&B project `acdc-cardiac-diffusion`.

4.2 Reconstruction Quality

Table 1 reports Chamfer distance and Hausdorff distance on the LV validation set. Our SE(3)-equivariant model achieves substantially lower error than the PCA baseline, confirming that non-linear deep generative modelling of mesh geometry is beneficial.

Table 1: Reconstruction quality on the LV validation set. Lower is better. $\pm$ denotes standard deviation across subjects. † results pending full baseline training runs.

Model	Chamfer (mm) $\downarrow$	Hausdorff (mm) $\downarrow$
PCA shape model (B1)	8.40 ± 1.04	32.43 ± 1.12
Voxel diffusion (B2)	**[TODO:]**†	**[TODO:]**†
PointNet++ VAE (B3)	**[TODO:]**†	**[TODO:]**†
MLP flow matching (B4)	**[TODO:]**†	**[TODO:]**†
Ours (EGNN + CFM)	$\mathbf{0.968 \pm 0.078}$	$\mathbf{1.576 \pm 0.094}$

4.3 Clinical Validity

Table 2 reports ejection fraction (EF) error and myocardial volume conservation for generated shapes compared to ground-truth ES meshes. The high EF error reflects that the model has not yet fully converged (200 epochs of training); we expect these metrics to improve substantially with longer training.

Table 2: Clinical validity metrics for the main model. EF error is mean absolute percentage error against ground-truth ejection fraction. Myo vol change is the mean percentage change in myocardial volume between generated ED and ES shapes. † results pending full baseline training runs.

Model	EF error (%) $\downarrow$	EF $\leq 5\%$ (%) $\uparrow$	Myo vol change (%) $\downarrow$
PCA (B1)	**[TODO:]**†	**[TODO:]**†	**[TODO:]**†
MLP flow (B4)	**[TODO:]**†	**[TODO:]**†	**[TODO:]**†
Ours	229.8 ± 570.8	3.0	881.0 ± 4612.7

4.4 Distribution Quality

Table 3 reports 1-nearest-neighbour accuracy (1-NNA) and linear pathology probe accuracy on the latent space. A 1-NNA of 0.595 indicates our model generates shapes that are partially indistinguishable from real ones, with room for improvement as training converges. The pathology probe accuracy of 32% is above chance (20% for five classes), showing the latent space encodes pathology-relevant structure even without label supervision.

Table 3: Distribution quality. 1-NNA closer to 0.5 is better (0.5 = perfect, 1.0 = mode collapse). Pathology accuracy is from a linear probe on the latent codes, trained with cross-validation. † results pending full baseline training runs.

Model	1-NNA $\rightarrow 0.5$	Pathology acc. (%) $\uparrow$
PCA (B1)	**[TODO:]**†	**[TODO:]**†
MLP flow (B4)	**[TODO:]**†	**[TODO:]**†
Ours	0.600	33.0

4.5 Ablation: Equivariance Under Rotation

[TODO: Show performance of main model vs. B4 as a function of training set size (e.g., 10%, 25%, 50%, 100% of training data). Expected finding: equivariant model degrades more gracefully in low-data regime.]

4.6 Hyperparameter Ablations

We ablate four design dimensions of the model, each starting from the baseline configuration (4 flow layers, latent dim 64, $\beta = 10^{-3}$, no KL warmup) and varying one or two factors at a time. Results are reported in Table 4.

KL weight β. Increasing β from 10^{-3} to 5×10^{-3} improves generation quality (1-NNA from 0.605 to 0.590) at a modest reconstruction cost (+0.08 mm Chamfer). Doubling further to 10^{-2} yields no additional gain in 1-NNA while further degrading reconstruction, suggesting 5×10^{-3} is near-optimal.

KL warmup. Linearly ramping β from 0 over the first 50 epochs (A2) recovers most of the reconstruction quality lost by the higher KL weight, reaching Chamfer 1.011 mm vs. 1.072 mm without warmup, while retaining the generation improvement (1-NNA 0.595).

Latent dimension. Doubling the latent dimension to 128 (A3) improves pathology probe accuracy from 32% to 35%, confirming the larger code captures more disease-relevant variation. Reconstruction and 1-NNA are largely unaffected compared to the baseline, suggesting the bottleneck is not latent capacity per se, but rather the degree of VAE regularisation.

Flow model depth. Increasing the velocity field from 4 to 6 EGNN layers (A4) yields the best reconstruction (Chamfer 0.967 mm, -0.027 mm vs. baseline) but slightly hurts 1-NNA ($+0.005$), indicating a more expressive velocity field sharpens shape fidelity at the cost of minor distributional degradation.

Best configuration. Combining KL warmup with 6 flow layers (A2+A4) achieves the best overall trade-off: Chamfer 0.968 ± 0.078 mm, Hausdorff 1.576 ± 0.094 mm, 1-NNA 0.600, and pathology probe 33%. This is the configuration used for all main results.

5 Discussion

[TODO: 3–4 paragraphs covering:

- **What the results say about the value of SE(3) equivariance.**

Table 4: Hyperparameter ablations on the LV validation set. Baseline uses 4 flow layers, latent dim 64, β=10^{-3}, no warmup. Best per-column values are **bold**. † Combined A2+A4 is the configuration used as "Ours".

Configuration	Chamfer↓	Hausdorff↓	1-NNA→0.5	Probe (%)↑
Baseline	0.995	1.614	0.605	32
A1: β=5×10^{-3}	1.072	1.765	**0.590**	33
A2: β=5×10^{-3} + warmup	1.011	1.646	0.595	32
A3: latent dim 128	1.018	1.655	0.595	**35**
A4: 6 flow layers	0.967	**1.566**	0.610	30
A3+A4: latent 128 + 6 layers	**0.963**	1.576	0.610	**35**
β=10^{-2} + warmup	1.042	1.697	0.595	31
A2+A4 (Ours)†: warmup + 6 layers	0.968	1.576	0.600	33

- **Limitations: single-structure, ED/ES only (not full cycle), Procrustes registration assumes rigid alignment.**
- **Future work: multi-structure joint model, 4D temporal modelling, learned registration.**

]

6 Conclusion

[TODO: 4–6 sentences. Restate problem, method, and three key findings.]

References

M. S. Albergo and E. Vanden-Eijnden. Building normalizing flows with stochastic interpolants. *International Conference on Learning Representations (ICLR)*, 2023.

O. Bernard, A. Lalande, C. Zotti, F. Cervenansky, X. Yang, P.-A. Heng, I. Cetin, K. Lekadir, O. Camara, M. A. G. Ballester, et al. Deep learning techniques for automatic MRI cardiac multi-structures segmentation and diagnosis: Is the problem solved? *IEEE Transactions on Medical Imaging*, 37(11):2514–2525, 2018.

T. F. Cootes, C. J. Taylor, D. H. Cooper, and J. Graham. Active shape models—their training and application. *Computer Vision and Image Understanding*, 61(1):38–59, 1995.

Y. Lipman, R. T. Q. Chen, H. Ben-Hamu, M. Nickel, and M. Le. Flow matching for generative modeling. *International Conference on Learning Representations (ICLR)*, 2023.

W. E. Lorensen and H. E. Cline. Marching cubes: A high resolution 3D surface construction algorithm. *ACM SIGGRAPH Computer Graphics*, 21(4):163–169, 1987.

V. G. Satorras, E. Hoogeboom, and M. Welling. E(n) equivariant graph neural networks. *International Conference on Machine Learning (ICML)*, 2021.

Y. Song, J. Sohl-Dickstein, D. P. Kingma, A. Kumar, S. Ermon, and B. Poole. Score-based generative modeling through stochastic differential equations. *International Conference on Learning Representations (ICLR)*, 2021.